Design-Partner Outreach Shortlist

Living GTM working doc (not product canon), derived from the dual-track research (archived at `archive/2026-07-12-dual-track-design-research.md`). This is a plan for landing the first design partners — the hands-on first customers who use Arteamis unfinished and give constant feedback in exchange for white-glove concierge setup.

The strategy: land 2 first partners, from 2 categories

Not five categories in parallel — a tiny team doing hands-on concierge can't. Land **two**, chosen to answer two different questions:

1. **One Web3 studio (via Morca's network)** — the *fastest door*. Warm intro, quick yes, gives you a reference and a real first deployment. Tests: can we run the concierge motion and does the loop deliver value?
2. **One AI dev agency (Category B)** — the *deepest pain*. Agencies rebuild context *per client engagement*, which is acute, repeated, and directly billable. Tests the sharper hypothesis: is the pain and willingness-to-pay actually higher here?

Web3 gives speed; the agency tests whether the durable wedge is somewhere else. If the agency partner's pain is clearly sharper, that reshapes the whole GTM — which is exactly what you want to learn early.

Target categories, ranked

Rank	Category	Pain	Concierge fit	Willingness to pay	Role
1	**AI dev agencies / AI-native services**	High (per-client rebuild)	High (they sell hands-on)	High (billable leverage)	**Primary — deepest, most repeatable pain**
2	**AI-native vertical-agent product studios**	High (existential)	High (grasp "context pack" instantly)	Med–High	Parallel — highest technical fit, harder to reach cold
3	**Web3 / infra studios**	Medium	High (relationship-driven)	Medium (volatile)	**Fastest door — land partner #1 here, don't over-index**
4	Indie multi-agent studios	Med–High (trace/WF-5)	Medium (often <5 people)	Low–Med	Learning cohort — harden trace, not prove WTP
5	Fintech / RWA / DeFi teams	Med–High (audit value)	Medium	High (compliance-driven)	Defer — slow sales cycle, revisit post-validation

Named leads to *qualify* (not validated targets)

The research surfaced these; their heavy agent use is **self-reported or inferred, not verified.** Treat them as first outreach + qualification, not confirmed fits.

- **AI dev agencies (Cat B):** Zensite (SG/KL), Codigo (SG), Seven Peaks (Bangkok), Flowtrix (Bengaluru) — SEA/India, within ~5h of Vietnam.
- **Vertical-agent studios (Cat A):** Context Studios (AI-native, agents in prod); the Anthropic "Claude Build Day" Seoul cohort; reach via Cursor's SG/JP/India developer communities.
- **Web3 / home turf (Cat C):** Pando Infinity, Cyberk, Wakumo Vietnam, Golden Sea Studio — plus your own direct SEA/Vietnam Web3 network (the warmest intros).

Qualify before you invest concierge time

A lead is a real candidate only if:

- [ ] They use **Claude Code / Cursor / MCP daily** (non-negotiable).
- [ ] They do **repeated client delivery** *or* ship agent-heavy products fast.
- [ ] **5–50 people** (below 5 often can't sustain the loop; above 50 is a different sale).
- [ ] The pain is **acute**, not just present — they can name the last time an agent got their context wrong and what it cost.

The ask (the concierge offer)

> "We're building this. Use it free for 30 days. We'll **hand-build your first Studio Brain** from one live client engagement / your codebase — you don't start from an empty brain. In exchange: **daily use, weekly feedback, and two studio intros.**"

The hand-built starter (≈5–15 team rules covering your most-repeated agent context — how you deploy, code conventions, "how X is handled" — plus enough sources to answer week-one questions) is what makes value show up on day one instead of week six. It's also the direct fix for the "slow time-to-value" risk.

Outreach angles per category

- **AI agencies (primary):** "*Every new client is a new hire starting from zero. Stop rebuilding context per engagement.*" Lead with the proof point that this is real ("3× client capacity, same team" — a competitor's own testimonial).
- **Vertical-agent studios (parallel):** "*Stop writing agent context from scratch every morning.*" Lead with the engineer's daily 11-minute context rebuild, plus the trace + fail-closed security story.

- **Web3 (partner #1):** warm, honest — "*We're building this, use it free for 30 days, tell us the truth, intro us to two studios.*" Caveat: confirm the pain is *acute*, not merely present — some Web3 teams are too code-centric to feel the knowledge-governance pain.

Selling points (for studio & investor conversations)

1. **"Git for your team's knowledge and agent context."** The one-liner.
2. **Trustworthy, not just shared** — personal-first governance; nothing becomes team truth (or reaches an agent) without a human review. The closest competitor can't say this.
3. **Trace with "why," not just a log** — every agent action is a governed record.
4. **Fail-closed security** — all 5 abuse categories tested and blocked.

Competitor note — don't panic, do move

- The **$50M "Inherent" is a *different company*** (London "Inherent Labs," ex-DeepMind). The real competitor, **inherent.sh**, is a single-person Indian company, no funding, one testimonial, launched last quarter.
- **inherent.sh auto-ingests everything via connectors — no personal→promoted→ governed arc.** No Personal Brain, no human-reviewed promotion. That gap is our differentiator.
- **Window is now:** YC's Summer 2026 RFS (Tom Blomfield) is near-verbatim our pitch — a tailwind, and a signal more competition is coming. Speed to real design partners beats more building.

Pricing — decided LAST, after validation

Do **not** charge day one. Convert at ~month 2 on a value trigger (weekly traceable agent output + successful concierge handoff to their team). Target **$200–500/mo** for a 5–50-person studio, **priced on governed work units** (context packs delivered, loops completed, trace-checked outputs) — *not* per seat (small teams resist per-seat "it adds up fast"). Anchors: Dust sells $30–150/seat/mo. These are directional; lock pricing only after the first partners run real usage.

Caveats

- Named studios' agent use is self-reported/inferred — qualify, don't assume.
- The 5–15-rule "minimum viable Studio Brain" is a hypothesis to measure in the first concierge deployment, not an established figure.
- Market tailwinds (YC RFS, vertical-AI enthusiasm, Gartner projections) are 2026 optimism — treat as claims, not facts.